

Homework 5

Sarah Cannon

2024-10-01

Classmates/other resources consulted: N/A

Be sure to load the tidyverse.

```
library(tidyverse)
```

Question 1 (6 points)

Parse the following dates and date/time combinations

a.

```
d_a <- "Sep. 25, (2024)"  
parse_datetime(d_a, "%b. %d, (%Y)")
```

```
## [1] "2024-09-25 UTC"
```

b.

```
d_b <- "2024-februar-12"  
parse_datetime(d_b, "%Y-%B-%d", locale = locale(date_names = "de"))
```

```
## [1] "2024-02-12 UTC"
```

(Hint: the language here is German; a list of the ISO 639-1 language abbreviations that R uses can be found at https://en.wikipedia.org/wiki/List_of_ISO_639-1_codes (https://en.wikipedia.org/wiki/List_of_ISO_639-1_codes))

c.

```
dt_c <- "February 13, 2023 at 7:45 am"  
parse_datetime(dt_c, "%B %d, %Y at %I:%M %p")
```

```
## [1] "2023-02-13 07:45:00 UTC"
```

Question 2 (9 points)

Consider the `dates_times.csv` file, in which the first column has a date and time, the second column has a date, and the third column has a time. Import this file, then parse all three columns so that they have the correct data types. You can either create new columns with the correct data types or replace the existing columns, it's up to you. Make sure the columns with the correct data types are the first three columns in your resulting tibble.

```
dates_times <- read_csv("dates_times.csv")
```

```
## Rows: 40 Columns: 3
## — Column specification —————
## Delimiter: ","
## chr (3): Date_times, Dates, Times
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
dates_times <- dates_times %>%
  mutate(
    Date_times_parsed = parse_datetime(Date_times, "%m-%d-%y: at %H:%M"),
    Dates_parsed = parse_date(Dates, "%d %B %Y", locale = locale(date_names = "fr")),
    Times_parsed = parse_time(Times, "%M minutes after %H %p")
  ) %>%
  select(Date_times_parsed, Dates_parsed, Times_parsed, everything())
dates_times
```

```
## # A tibble: 40 × 6
##   Date_times_parsed  Dates_parsed Times_parsed Date_times      Dates      Time
##   <dtm>            <date>      <time>      <chr>          <chr>      <chr>
>
## 1 2021-08-01 16:43:00 2021-04-01 16:34      8-1-21: at 16:43 1 Avril 2021 34 m
inutes after 4pm
## 2 2021-08-02 17:16:00 2021-04-02 13:35      8-2-21: at 17:16 2 Avril 2021 35 m
inutes after 1pm
## 3 2021-08-03 16:32:00 2021-04-03 16:34      8-3-21: at 16:32 3 Avril 2021 34 m
inutes after 4pm
## 4 2021-08-04 16:32:00 2021-04-04 18:54      8-4-21: at 16:32 4 Avril 2021 54 m
inutes after 6pm
## 5 2021-08-05 16:23:00 2021-04-05 16:23      8-5-21: at 16:23 5 Avril 2021 23 m
inutes after 4pm
## 6 2021-08-06 17:12:00 2021-04-06 17:23      8-6-21: at 17:12 6 Avril 2021 23 m
inutes after 5pm
## 7 2021-08-07 16:55:00 2021-04-07 16:25      8-7-21: at 16:55 7 Avril 2021 25 m
inutes after 4pm
## 8 2021-08-08 16:35:00 2021-04-08 21:45      8-8-21: at 16:35 8 Avril 2021 45 m
inutes after 9pm
## 9 2021-08-09 17:01:00 2021-04-09 16:25      8-9-21: at 17:01 9 Avril 2021 25 m
inutes after 4pm
## 10 2021-08-10 17:49:00 2021-04-10 16:12      8-10-21: at 17:49 10 Avril 2021 12 m
inutes after 4pm
## # i 30 more rows
```

Question 3 (4 points)

Import the attached data set “Monthly_amounts.txt”. The first column of this data set contains year and month information; parse it so that it is of the correct data type. While the original data has no day information, what happens in your new column?

```
monthly_amounts <- read_csv("Monthly_Amounts.txt", skip = 1)
```

```
## Warning: One or more parsing issues, call `problems()` on your data frame for details, e.g.:
##   dat <- vroom(...)
##   problems(dat)
```

```
## Rows: 20 Columns: 3
## — Column specification —————
## Delimiter: ","
## chr (3): Year_month, largest_amount, average_amount
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
mutate(monthly_amounts, Year_month_parsed = parse_date(Year_month, "%Y-%m"))
```

```
## # A tibble: 20 × 4
##   Year_month largest_amount average_amount Year_month_parsed
##   <chr>      <chr>          <chr>          <date>
## 1 2023-1      63            45.7          2023-01-01
## 2 2023-2      56            35.2          2023-02-01
## 3 2023-3      54            34.2          2023-03-01
## 4 2023-4      12            11.6          2023-04-01
## 5 2023-5      <NA>          35.2          2023-05-01
## 6 2023-6      87            56.9          2023-06-01
## 7 2023-7      82            45.7          2023-07-01
## 8 2023-8      36            17.9          2023-08-01
## 9 2023-9      98            54.3          2023-09-01
## 10 2023-10     16            9.8           2023-10-01
## 11 2023-11     78            50.2          2023-11-01
## 12 2023-12     45            43.6          2023-12-01
## 13 2024-1      35            31.2          2024-01-01
## 14 2024-2      91            -             2024-02-01
## 15 2024-3      45            34.5          2024-03-01
## 16 2024-4      67            34.6          2024-04-01
## 17 2024-5      34            21.9          2024-05-01
## 18 2024-6      *             49.8          2024-06-01
## 19 2024-7      35            23.8          2024-07-01
## 20 2024-8      45            43            2024-08-01
```

The new column will have the first day of the month as the default day. Assuming we can ignore the warning coming from the extra column in one row.

Question 4 (6 points)

- Create a string in R containing the following sentence, including its punctuation: It's sunny today, but he said, "It'll be rainy tomorrow." To be sure you've made the correct string, print it out using the `writeLines()` function.

```
sentence <- "It's sunny today, but he said, \"It'll be rainy tomorrow.\""  
writeLines(sentence)
```

```
## It's sunny today, but he said, "It'll be rainy tomorrow."
```

- b. Explain the difference between the strings: “a b” and “a n b”. The answer is not just that one has an extra space and one doesn’t. Your explanation should mention escape characters. Note: to make a backslash show up in your explanation in your knitted file, you must use two backslashes, like \\. Be sure to check your knitted file to make sure all parts of your explanation are appearing correctly.**

The string “a \n b” contains an escape character `\n`, which represents a newline. Therefore, when printed, it will display as:

a

b

On the other hand, “a \ n b” does not contain a valid escape character due to the space after the backslash, so it will be printed as is, with a space and a backslash:

a \ n b

Question 5 (4 points)

This question references the following strings

```
s1 <- "the cat, gracie, is sleepy"
```

```
s2 <- "The Dog Is Sleepy Too!"
```

- a. Make s1 uppercase**

```
s1_upper <- str_to_upper(s1)  
s1_upper
```

```
## [1] "THE CAT, GRACIE, IS SLEEPY"
```

- b. Make s2 lowercase**

```
s2_lower <- str_to_lower(s2)
s2_lower
```

```
## [1] "the dog is sleepy too!"
```

c. Make the first letter of every word in s1 capitalized, while all other letters are lowercase.

```
s1_title <- str_to_title(s1)
s1_title
```

```
## [1] "The Cat, Gracie, Is Sleepy"
```

d. Write a command that will output the number of characters in string s2.

```
s2_length <- str_length(s2)
s2_length
```

```
## [1] 22
```

Question 6 (6 points)

In the U.S., mailing addresses have zipcodes consisting of five digits, then a dash, then four digits. An example might be 91711-4285. Suppose you have a tibble, like the following example, where the first five digits are in a different column than the last four digits.

```
zip_codes <- tibble(Zip = c("91711-3452", "20322-3009", "93782-8473", "78392-8762", "87639-2563", "47628-5416", "20874-5726"))
```

a. Use the `str_sub()` function to split the Zip column into two columns, one with the first five digits of the zip code and one with the last four digits of the zip code.

```
zip_codes <- zip_codes %>%
  mutate(
    first_five = str_sub(Zip, 1, 5),
    last_four = str_sub(Zip, 7, 10)
  )
zip_codes
```

```
## # A tibble: 7 × 3
##   Zip      first_five last_four
##   <chr>      <chr>      <chr>
## 1 91711-3452 91711      3452
## 2 20322-3009 20322      3009
## 3 93782-8473 93782      8473
## 4 78392-8762 78392      8762
## 5 87639-2563 87639      2563
## 6 47628-5416 47628      5416
## 7 20874-5726 20874      5726
```

b. Use the `separate()` function to split the `Zip` column into two columns, one with the first five digits of the zip code and one with the last four digits of the zip code.

```
zip_codes <- zip_codes %>%
  separate(Zip, into = c("first_five", "last_four"), sep = "-")
zip_codes
```

```
## # A tibble: 7 × 2
##   first_five last_four
##   <chr>      <chr>
## 1 91711      3452
## 2 20322      3009
## 3 93782      8473
## 4 78392      8762
## 5 87639      2563
## 6 47628      5416
## 7 20874      5726
```

Question 7 (8 points)

At a particular company, an employee's email address consists of their first initial, their middle initial (if they have one), their last name, and the last two digits of their Employee ID number, followed by "@company.com". For example, for an employee Alice A. Smith with employee ID number 45398545, her email address would be AASmith45@company.com (mailto:AASmith45@company.com). For an employee Bob Jones (who does not have a middle initial) with employee ID number 345582, his email address would be BJones82@company.com (mailto:BJones82@company.com). For the table below, write code to add a new column consisting of each employee's email address, computed from the values in the other columns.

```
employees <- tibble(
  FirstName = c("Alice", "Bob", "Simba", "Nala", "Timon", "Pumbaa", "Rafiki", "Scar"),
  MiddleInitial = c("A", NA, "E", "Q", "P", NA, "P", "L"),
  LastName = c("Smith", "Jones", "Clark", "Davis", "Evans", "Frank", "Ghosh", "Hills"),
  EmployeeID = c(45398545, 345582, 2354463, 345346, 2346377022, 20345423, 20223454, 204254))
employees
```

```
## # A tibble: 8 × 4
##   FirstName MiddleInitial LastName EmployeeID
##   <chr>      <chr>      <chr>      <dbl>
## 1 Alice      A          Smith      45398545
## 2 Bob        <NA>       Jones      345582
## 3 Simba      E          Clark      2354463
## 4 Nala       Q          Davis      345346
## 5 Timon      P          Evans      2346377022
## 6 Pumbaa     <NA>       Frank      20345423
## 7 Rafiki     P          Ghosh      20223454
## 8 Scar       L          Hills      204254
```

```
employees <- employees %>%
  mutate(
    Email = str_c(
      str_sub(FirstName, 1, 1),
      ifelse(is.na(MiddleInitial), "", MiddleInitial),
      LastName,
      str_sub(as.character(EmployeeID), -2, -1),
      "@company.com"
    )
  )
employees
```



```
## # A tibble: 8 × 5
##   FirstName MiddleInitial LastName EmployeeID Email
##   <chr>      <chr>      <chr>      <dbl> <chr>
## 1 Alice      A          Smith      45398545 AASmith45@company.com
## 2 Bob        <NA>       Jones      345582 BJones82@company.com
## 3 Simba      E          Clark      2354463 SEClark63@company.com
## 4 Nala       Q          Davis      345346 NQDavis46@company.com
## 5 Timon      P          Evans      2346377022 TPEvans22@company.com
## 6 Pumbaa     <NA>       Frank      20345423 PFrank23@company.com
## 7 Rafiki     P          Ghosh      20223454 RPGhosh54@company.com
## 8 Scar       L          Hills      204254 SLHills54@company.com
```

Question 8 (12 points)

In each part, say whether the data is tidy or not, and explain why.

a.

```
## # A tibble: 10 × 4
##   Team_Abbreviation Team_Name      Division `Wins-Losses`
##   <chr>            <chr>      <chr>      <chr>
## 1 TB              Tampa Bay Rays East      100-62
## 2 BOS             Boston Red Sox East      92-70
## 3 NYY             New York Yankees East      92-70
## 4 TOR             Toronto Blue Jays East      91-71
## 5 BAL             Baltimore Orioles East      52-110
## 6 CHW             Chicago White Sox Central    93-69
## 7 CLE             Cleveland Indians Central    80-82
## 8 DET             Detroit Tigers Central    77-85
## 9 KC              Kansas City Royals Central    74-88
## 10 MIN            Minnesota Twins Central    73-89
```

This data is not tidy because the `Wins-Losses` column contains two values (wins and losses) in a single entry. Each column should contain only one value.

b.

```
## # A tibble: 10 × 4
##   Name      College Info      Value
##   <chr>      <chr>  <chr>      <dbl>
## 1 Student A CMC      GPA         3.8
## 2 Student B CMC      GPA         3.7
## 3 Student C Pitzer   GPA         3.72
## 4 Student D CMC      GPA         3.66
## 5 Student E Scripps  GPA         3.72
## 6 Student A CMC      Graduation 2022
## 7 Student B CMC      Graduation 2024
## 8 Student C Pitzer   Graduation 2023
## 9 Student D CMC      Graduation 2023
## 10 Student E Scripps Graduation 2023
```

This data is not tidy because each student has multiple rows for different types of information (GPA, Graduation). Each row should represent a single observation (a student), and each column should represent a single variable.

C.

```
## # A tibble: 7 × 5
##   Day      Temperature_F Wind_mph UV_index ChanceOfRain_percent
##   <chr>          <dbl>    <dbl>    <dbl>          <dbl>
## 1 Thursday         71         7         2             70
## 2 Friday           63        12         4             80
## 3 Saturday         71        10         7              4
## 4 Sunday           78        11         7              0
## 5 Monday           71        13         6              7
## 6 Tuesday          70        11         6              0
## 7 Wednesday        74        10         6              0
```

This data is tidy because each row represents an observation, each column represents a variable, and each entry contains only one value.

d.

```
## # A tibble: 5 × 6
##   cut      extra_small small medium large extra_large
##   <chr>          <dbl> <dbl>  <dbl> <dbl>          <dbl>
## 1 Fair           32     23     34     23           34
## 2 Good           45     45     45     21           56
## 3 Very Good      67     26     63     43           23
## 4 Premium        32     78     78     47           14
## 5 Ideal          14     23     99     21           21
```

This data is tidy because each row represents an observation (a diamond cut), and each column represents a variable.

Question 9 (12 points)

For each of these tibbles, perform the necessary operation to make it tidy.

a.

```
## # A tibble: 12 × 3
##   month      metric      average
##   <chr>      <chr>      <dbl>
## 1 September high_temperature  89
## 2 September low_temperature   60
## 3 September rain_inches      0.15
## 4 September daylight_hours  12.5
## 5 October   high_temperature  80
## 6 October   low_temperature   55
## 7 October   rain_inches      1.05
## 8 October   daylight_hours  11.5
## 9 November  high_temperature  74
## 10 November low_temperature   47
## 11 November rain_inches      1.62
## 12 November daylight_hours  10.5
```

```
avg_weather_tidy <- avg_weather %>%
  pivot_wider(names_from = metric, values_from = average)
avg_weather_tidy
```

```
## # A tibble: 3 × 5
##   month      high_temperature low_temperature rain_inches daylight_hours
##   <chr>      <dbl>          <dbl>      <dbl>      <dbl>
## 1 September      89            60        0.15      12.5
## 2 October       80            55        1.05      11.5
## 3 November      74            47        1.62      10.5
```

b.

```
## # A tibble: 6 × 2
##   Chemical_Name Safe_Temperature_Range
##   <chr>          <chr>
## 1 Chemical 1    32-212
## 2 Chemical 2    50-100
## 3 Chemical 3    45-48
## 4 Chemical 4    40-345
## 5 Chemical 5    100-250
## 6 Chemical 6    112-140
```

```
chemicals_tidy <- chemicals %>%
  separate(Safe_Temperature_Range, into = c("Min_Temperature", "Max_Temperature"), sep =
    "_")
chemicals_tidy
```

```
## # A tibble: 6 × 3
##   Chemical_Name Min_Temperature Max_Temperature
##   <chr>         <chr>         <chr>
## 1 Chemical 1    32             212
## 2 Chemical 2    50             100
## 3 Chemical 3    45             48
## 4 Chemical 4    40             345
## 5 Chemical 5    100            250
## 6 Chemical 6    112            140
```

c. Information about the following data is available at

<https://github.com/rfordatascience/tidytuesday/blob/master/data/2023/2023-07-11/readme.md>
(<https://github.com/rfordatascience/tidytuesday/blob/master/data/2023/2023-07-11/readme.md>)

```
## Rows: 144 Columns: 19
## — Column specification —————
## Delimiter: ","
## dbl (19): Year, Jan, Feb, Mar, Apr, May, Jun, Jul, Aug, Sep, Oct, Nov, Dec, J-D, D-N,
## DJF, MAM, JJA, SON
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
## # A tibble: 144 × 13
##   Year Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 1880 -0.19 -0.25 -0.09 -0.17 -0.1 -0.21 -0.18 -0.11 -0.15 -0.24 -0.22 -0.18
## 2 1881 -0.2 -0.15 0.03 0.05 0.05 -0.19 0 -0.04 -0.16 -0.22 -0.19 -0.08
## 3 1882 0.16 0.13 0.04 -0.16 -0.14 -0.22 -0.17 -0.08 -0.15 -0.24 -0.17 -0.36
## 4 1883 -0.3 -0.37 -0.13 -0.19 -0.18 -0.08 -0.08 -0.14 -0.23 -0.12 -0.24 -0.11
## 5 1884 -0.13 -0.09 -0.37 -0.4 -0.34 -0.35 -0.31 -0.28 -0.28 -0.25 -0.34 -0.31
## 6 1885 -0.59 -0.34 -0.27 -0.42 -0.45 -0.44 -0.34 -0.32 -0.29 -0.24 -0.24 -0.11
## 7 1886 -0.44 -0.51 -0.43 -0.28 -0.24 -0.35 -0.18 -0.31 -0.24 -0.28 -0.28 -0.26
## 8 1887 -0.72 -0.57 -0.36 -0.35 -0.31 -0.25 -0.26 -0.36 -0.26 -0.36 -0.27 -0.33
## 9 1888 -0.34 -0.36 -0.41 -0.2 -0.22 -0.17 -0.11 -0.16 -0.12 0.01 0.03 -0.04
## 10 1889 -0.09 0.16 0.06 0.1 -0.01 -0.1 -0.08 -0.2 -0.24 -0.25 -0.33 -0.29
## # i 134 more rows
```

```
global_temps_tidy <- global_temps %>%
  pivot_longer(cols = Jan:Dec, names_to = "Month", values_to = "Temperature")
global_temps_tidy
```

```
## # A tibble: 1,728 × 3
##   Year Month Temperature
##   <dbl> <chr>      <dbl>
## 1  1880 Jan        -0.19
## 2  1880 Feb        -0.25
## 3  1880 Mar        -0.09
## 4  1880 Apr        -0.17
## 5  1880 May        -0.1
## 6  1880 Jun        -0.21
## 7  1880 Jul        -0.18
## 8  1880 Aug        -0.11
## 9  1880 Sep        -0.15
## 10 1880 Oct        -0.24
## # i 1,718 more rows
```

Question 10 (9 points)

a. Consider the following example table.

```
diamonds_counts <- tibble(cut = c("Fair", "Good", "Very Good", "Premium", "Ideal"),
  extra_small = c(32,45,67,32,14),
  small = c(23,45,26,78,23),
  medium = c(34,45,63,78,99),
  large = c(23,21,43,47,21),
  extra_large = c(34,56,23,14,21))
diamonds_counts
```

```
## # A tibble: 5 × 6
##   cut      extra_small small medium large extra_large
##   <chr>      <dbl> <dbl>  <dbl> <dbl>      <dbl>
## 1 Fair          32    23    34    23         34
## 2 Good          45    45    45    21         56
## 3 Very Good     67    26    63    43         23
## 4 Premium       32    78    78    47         14
## 5 Ideal         14    23    99    21         21
```

Explain in your own words why the following two code chunks produce the same tibble.

```
diamonds_counts %>%
  pivot_longer(extra_small:extra_large, names_to = "Size", values_to = "Count") %>%
  filter(Size %in% c("extra_small", "extra_large"))
```

```
## # A tibble: 10 × 3
##   cut      Size      Count
##   <chr>    <chr>    <dbl>
## 1 Fair    extra_small  32
## 2 Fair    extra_large  34
## 3 Good    extra_small  45
## 4 Good    extra_large  56
## 5 Very Good extra_small  67
## 6 Very Good extra_large  23
## 7 Premium extra_small  32
## 8 Premium extra_large  14
## 9 Ideal   extra_small  14
## 10 Ideal  extra_large  21
```

```
diamonds_counts %>%
  select(cut, extra_small, extra_large) %>%
  pivot_longer(extra_small:extra_large, names_to = "Size", values_to = "Count")
```

```
## # A tibble: 10 × 3
##   cut      Size      Count
##   <chr>    <chr>    <dbl>
## 1 Fair    extra_small  32
## 2 Fair    extra_large  34
## 3 Good    extra_small  45
## 4 Good    extra_large  56
## 5 Very Good extra_small  67
## 6 Very Good extra_large  23
## 7 Premium extra_small  32
## 8 Premium extra_large  14
## 9 Ideal   extra_small  14
## 10 Ideal  extra_large  21
```

Both code chunks produce the same tibble because they both transform the data from wide to long format, focusing only on the `extra_small` and `extra_large` columns. They just do it in a different order with the first code chunk filtering the data after it's pivoted and the second code chunk filtering/selecting the data before it's pivoted.

b. Why doesn't the following code work as expected? Explain what went wrong here, and why the `pivot_wider` function doesn't work for this data set in the same way we learned in class.

```
cats <- tribble(
  ~name,      ~names, ~values,
  "Gracie the Cat", "age",      6.5,
  "Gracie the Cat", "height_in", 14,
  "Gracie the Cat", "age",      5,
  "Patches the Cat", "age",      2,
  "Patches the Cat", "height_in", 11
)

newtibble <- cats %>% pivot_wider(names_from = names, values_from = values)
```

```
## Warning: Values from `values` are not uniquely identified; output will contain list-c
ols.
## • Use `values_fn = list` to suppress this warning.
## • Use `values_fn = {summary_fun}` to summarise duplicates.
## • Use the following dplyr code to identify duplicates.
## {data} |>
## dplyr::summarise(n = dplyr::n(), .by = c(name, names)) |>
## dplyr::filter(n > 1L)
```

```
newtibble
```

```
## # A tibble: 2 × 3
##   name      age      height_in
##   <chr>    <list>   <list>
## 1 Gracie the Cat <dbl [2]> <dbl [1]>
## 2 Patches the Cat <dbl [1]> <dbl [1]>
```

The code doesn't work because there are 2 `age` values for Gracie the Cat but the `pivot_wider` function requires unique combinations of the `name` and `names` columns.

c. Look up what the tidyverse's `spread()` and `gather()` functions do and explain them below. These functions are no longer under active development, but exist in a lot of previously written code. Which functions we've learned recently are the updated versions of `spread` and `gather`?

The `spread()` function in the tidyverse is used to turn key-value pairs into columns, effectively making the data wider. The `gather()` function is used to turn columns into key-value pairs, making the data longer. They functions have been obsoleted by `pivot_wider()` and `pivot_longer()`.

Question 11 (12 points)

Consider the following data set; more information is available at

<https://github.com/rfordatascience/tidytuesday/blob/master/data/2020/2020-03-24/readme.md>

(<https://github.com/rfordatascience/tidytuesday/blob/master/data/2020/2020-03-24/readme.md>)

```
brain_injuries <- read_csv("https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2020/2020-03-24/tbi_age.csv") %>% select(age_group:number_est) %>% filter(age_group != "Total")
```

```
## Rows: 231 Columns: 5
```

```
## — Column specification —————
```

```
## Delimiter: ","
```

```
## chr (3): age_group, type, injury_mechanism
```

```
## dbl (2): number_est, rate_est
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
brain_injuries
```



```
## # A tibble: 210 × 4
##   age_group type          injury_mechanism
##   <chr>      <chr>          <chr>
<dbl>
## 1 0-17      Emergency Department Visit Motor Vehicle Crashes
47138
## 2 0-17      Emergency Department Visit Unintentional Falls
397190
## 3 0-17      Emergency Department Visit Unintentionally struck by or against an object
229236
## 4 0-17      Emergency Department Visit Other unintentional injury, mechanism unspecified
55785
## 5 0-17      Emergency Department Visit Intentional self-harm
NA
## 6 0-17      Emergency Department Visit Assault
24360
## 7 0-17      Emergency Department Visit Other or no mechanism specified
57983
## 8 0-4       Emergency Department Visit Motor Vehicle Crashes
5464
## 9 0-4       Emergency Department Visit Unintentional Falls
230776
## 10 0-4      Emergency Department Visit Unintentionally struck by or against an object
53436
## # i 200 more rows
```

a. (3 points) **Make a table that displays all the different values in the `age_group` category and how many times each appears in the data set.**

```
brain_injuries %>%
  group_by(age_group) %>%
  summarize(count = n())
```

```
## # A tibble: 10 × 2
##   age_group count
##   <chr>      <int>
## 1 0-17        21
## 2 0-4         21
## 3 15-24       21
## 4 25-34       21
## 5 35-44       21
## 6 45-54       21
## 7 5-14        21
## 8 55-64       21
## 9 65-74       21
## 10 75+        21
```

- b. (3 points) **Some of these age groups are overlapping, which is not ideal for data analysis. For example, it's currently very challenging to look in-depth at the age groups 15-17 and 18-24. To simplify the age groups, it's actually easiest if they're columns! Use a function we learned this week to turn the age_group values into columns, where the values in those columns come from the number_est column.**

```
brain_injuries %>%
  pivot_wider(names_from = age_group, values_from = number_est)
```

```
## # A tibble: 21 × 12
##   type                                injury_mechanism      `0-17`  `0-4`  `5-14`  `15
-24`  `25-34`  `35-44`  `45-54`  `55-64`  `65-74`  `75+`
##   <chr>                                <chr>                <dbl>  <dbl>  <dbl>  <
dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
## 1 Emergency Department Visit Motor Vehicle Crashes      47138   5464   19785   10
3892   71641   44108   40020   27193   13829   8176
## 2 Emergency Department Visit Unintentional Falls          397190  230776  133084    9
6568   70210   68830   95127  112460  120327  286031
## 3 Emergency Department Visit Unintentionally struck by or ... 229236   53436  120839   10
6679   44404   32479   30495   20408   11937   13270
## 4 Emergency Department Visit Other unintentional injury, m... 55785   12007   30656    3
7118   22360   17541   17808   12928    7077   7440
## 5 Emergency Department Visit Intentional self-harm           NA      NA      NA
870     650     421     247     105     NA     NA
## 6 Emergency Department Visit Assault          24360     674   9690    6
5399   57213   34100   27682   11538    2893   1260
## 7 Emergency Department Visit Other or no mechanism specifi... 57983   19360   26022    3
3395   20974   16503   15962   13387   10051   17318
## 8 Hospitalizations      Motor Vehicle Crashes      5830     870   2395    1
2925   11050    7305    8490    7280    4485   3965
## 9 Hospitalizations      Unintentional Falls          7935   4700   2270
3910    4470    5640   12010   18490   25235   74005
## 10 Hospitalizations     Unintentionally struck by or ... 1985     510    980
1070     635     610     685     765     790   1045
## # i 11 more rows
```

- c. (3 points) **By adding and subtracting columns from the table you produced in the previous part, make new columns for ages 15-17 and 18-24. Hint: the number_est for the 15-17 age group is the estimate for 0-17 minus the estimate for 0-4 and the estimate for 5-14. Keep only the type column, injury_mechanism column, and the columns for the age groups 0-4, 5-14, 15-17, and 18-24.**

```
brain_injuries_wide <- brain_injuries %>%
  pivot_wider(names_from = age_group, values_from = number_est) %>%
  mutate(
    `15-17` = `0-17` - `0-4` - `5-14`,
    `18-24` = `0-4` + `5-14` + `15-24` - `0-17`
  ) %>%
  select(type, injury_mechanism, `0-4`, `5-14`, `15-17`, `18-24`)
brain_injuries_wide
```

```
## # A tibble: 21 × 6
##   type                injury_mechanism      `0-4`
##   `5-14` `15-17` `18-24`
##   <chr>      <chr>
> <dbl> <dbl> <dbl>
## 1 Emergency Department Visit Motor Vehicle Crashes      546
4 19785    21889    82003
## 2 Emergency Department Visit Unintentional Falls      23077
6 133084    33330    63238
## 3 Emergency Department Visit Unintentionally struck by or against an object    5343
6 120839    54961    51718
## 4 Emergency Department Visit Other unintentional injury, mechanism unspecified 1200
7 30656    13122    23996
## 5 Emergency Department Visit Intentional self-harm      N
A    NA      NA      NA
## 6 Emergency Department Visit Assault      67
4 9690    13996    51403
## 7 Emergency Department Visit Other or no mechanism specified    1936
0 26022    12601    20794
## 8 Hospitalizations      Motor Vehicle Crashes      87
0 2395    2565    10360
## 9 Hospitalizations      Unintentional Falls      470
0 2270     965    2945
## 10 Hospitalizations      Unintentionally struck by or against an object    51
0 980     495    575
## # i 11 more rows
```

d. (3 points) This data is not tidy! Use an appropriate function we learned this week to make the data tidy.

```
brain_injuries_tidy <- brain_injuries_wide %>%
  pivot_longer(cols = `0-4`: `18-24`, names_to = "age_group", values_to = "number_est")
brain_injuries_tidy
```

```
## # A tibble: 84 × 4
##   type                injury_mechanism      age_grou
p number_est
##   <chr>                <chr>            <chr>
<dbl>
## 1 Emergency Department Visit Motor Vehicle Crashes 0-4
5464
## 2 Emergency Department Visit Motor Vehicle Crashes 5-14
19785
## 3 Emergency Department Visit Motor Vehicle Crashes 15-17
21889
## 4 Emergency Department Visit Motor Vehicle Crashes 18-24
82003
## 5 Emergency Department Visit Unintentional Falls 0-4
230776
## 6 Emergency Department Visit Unintentional Falls 5-14
133084
## 7 Emergency Department Visit Unintentional Falls 15-17
33330
## 8 Emergency Department Visit Unintentional Falls 18-24
63238
## 9 Emergency Department Visit Unintentionally struck by or against an object 0-4
53436
## 10 Emergency Department Visit Unintentionally struck by or against an object 5-14
120839
## # i 74 more rows
```

Question 12 (8 points)

Consider the file GPAs.csv. Import this data set, and carefully make it both clean and tidy. There will be several steps involved in this process. (the GPAs were randomly generated and do not reflect actual student grades).

```
gpa <- read_csv("GPAs.csv")
```

```
## Rows: 19 Columns: 5
## — Column specification —————
## Delimiter: ","
## chr (2): College, Class
## dbl (3): Math_GPA, Data Science_GPA, Other_GPA
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
gpa
```

```
## # A tibble: 19 × 5
##   College Class      Math_GPA `Data Science_GPA` Other_GPA
##   <chr>   <chr>      <dbl>          <dbl>      <dbl>
## 1 CMC     Freshman    3.74           3.84       3.78
## 2 CMC     Sophomore   3.09           3.14       3.16
## 3 CMC     Junior      3.74           3.84       3.94
## 4 CMC     Senior      3.63           3.96       3.4
## 5 HMC     freshman    3.26           3.11       3.79
## 6 HMC     sophomore   3.58           3.78       3.64
## 7 HMC     senior      3.93           3.02       3.02
## 8 Pitzer  Freshman    3.33           3.46       3.96
## 9 Pitzer  Sophomore   3.01           3.93       3.9
## 10 Pitzer Junior      3.87           3.74       3.19
## 11 Pitzer Senior     3.25           3.41       3.23
## 12 Pomona freshman    3.7            3.7       3.32
## 13 Pomona sophomore   3.4            3.31       3.9
## 14 Pomona junior      3.24            3.82       3.25
## 15 Pomona senior      3.23            3.72       3.51
## 16 Scripps Freshman    3.6            3.75       3.04
## 17 Scripps Sophomore   3.51            3.34       3.14
## 18 Scripps Junior     3.68            3.45       3.2
## 19 Scripps Senior     3.15            3.53       3.53
```

```
gpa <- read_csv("GPAs.csv")
```

```
## Rows: 19 Columns: 5
## — Column specification —————
## Delimiter: ","
## chr (2): College, Class
## dbl (3): Math_GPA, Data Science_GPA, Other_GPA
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
gpa <- gpa %>%
  rename(Data_Science_GPA = `Data Science_GPA`) %>%
  mutate(Class = str_to_title(Class)) %>%
  complete(College, Class, fill = list(Grade = NA)) %>%
  pivot_longer(cols = Math_GPA:Other_GPA, names_to = "Type", values_to = "GPA") %>%
  mutate(Type = str_replace(Type, "_GPA", ""))
gpa
```

```
## # A tibble: 60 × 4
##   College Class      Type      GPA
##   <chr>      <chr>      <chr>      <dbl>
## 1 CMC        Freshman  Math        3.74
## 2 CMC        Freshman  Data_Science 3.84
## 3 CMC        Freshman  Other        3.78
## 4 CMC        Junior    Math        3.74
## 5 CMC        Junior    Data_Science 3.84
## 6 CMC        Junior    Other        3.94
## 7 CMC        Senior    Math        3.63
## 8 CMC        Senior    Data_Science 3.96
## 9 CMC        Senior    Other        3.4
## 10 CMC       Sophomore Math        3.09
## # i 50 more rows
```

Question 13 (4 points)

- a. (2 points) **Reflect on how this class is going for you so far - What's working well for you? What are you finding challenging? What changes can you make that might help improve your learning?**

I'm understanding the concepts of the course which is good. Doing the homework quickly is a challenge. I think it's more an issue of study habits and time management since I tend to procrastinate.

- b. (2 points) **With the first test coming up, how do you plan to study? What methods of studying do you think will be most effective? Outline a study plan here.**

I plan to study by systematically reviewing the class notes and activities. I will also collaborate with my homework group to review key concepts and help each other understand the material.