

Homework 6

Prof. Sarah Cannon

Due Thursday 10/17/2024

Question 1 (4 points)

Read the ACM Code of Ethics at <https://www.acm.org/code-of-ethics> (<https://www.acm.org/code-of-ethics>) in its entirety. In 1-2 paragraphs below, summarize which parts of the code of ethics are most relevant for a data scientist and how you plan to ensure you adhere to this code of ethics in any future data science (or data science-adjacent) work you do.

The ACM Code of Ethics is important for data scientists as it emphasizes principles like contributing to society, avoiding harm, and respecting privacy. For data scientists, these principles translate into ensuring that data collection and analysis benefit society, do not cause harm, and protect individuals' privacy. I plan to implement robust data privacy measures, ensure transparency in data usage, and prioritize ethical considerations in all data-driven decisions. Additionally, I will engage in continuous learning to stay updated on best practices and ethical standards in data science.

Question 2 (4 points)

Consider the sets $A = \{\text{red, green, blue}\}$ and $B = \{\text{purple, blue, pink, yellow}\}$.

a. For these sets A and B, what is the union of A and B?

{red, green, blue, purple, pink, yellow}

b. For these sets A and B, what is the intersection of A and B?

{blue}

c. For these sets A and B, what is set difference $A - B$?

{red, green}

d. For these sets A and B, what is the Cartesian Product $A \times B$?

{(red, purple), (red, blue), (red, pink), (red, yellow), (green, purple), (green, blue), (green, pink), (green, yellow), (blue, purple), (blue, pink), (blue, yellow)}

Question 3 (8 points)

This question deals with the Lahman package, which has several tables related to baseball. **MAKE SURE NO MORE THAN 10 ROWS OF ANY LIST OR TABLE PRINT IN YOUR KNITTED FILE.**

- a. **What column makes a primary key in the People table? Explain how you know this is a valid key.**

```
head(People, 10)
```

##	playerID	birthYear	birthMonth	birthDay	birthCity	birthCountry	birthState	death
Year	deathMonth	deathDay	deathCountry	deathState				
## 1	aardsda01	1981	12	27	Denver	USA	CO	
NA	NA	NA	<NA>	<NA>				
## 2	aaronha01	1934	2	5	Mobile	USA	AL	
2021	1	22	USA	GA				
## 3	aaronto01	1939	8	5	Mobile	USA	AL	
1984	8	16	USA	GA				
## 4	aasedo01	1954	9	8	Orange	USA	CA	
NA	NA	NA	<NA>	<NA>				
## 5	abadan01	1972	8	25	Palm Beach	USA	FL	
NA	NA	NA	<NA>	<NA>				
## 6	abadfe01	1985	12	17	La Romana	D.R.	La Romana	
NA	NA	NA	<NA>	<NA>				
## 7	abadijo01	1850	11	4	Philadelphia	USA	PA	
1905	5	17	USA	NJ				
## 8	abbated01	1877	4	15	Latrobe	USA	PA	
1957	1	6	USA	FL				
## 9	abbeybe01	1869	11	11	Essex	USA	VT	
1962	6	11	USA	VT				
## 10	abbeych01	1866	10	14	Falls City	USA	NE	
1926	4	27	USA	CA				
##	deathCity	nameFirst	nameLast	nameGiven	weight	height	bats	throws
debut	bbrefID	finalGame	retroID	deathDate				
## 1		<NA>	David Aardsma	David Allan	215	75	R	R 2
004-04-06	aardsda01	2015-08-23	aardd001	<NA>				
## 2	Atlanta	Hank	Aaron	Henry Louis	180	72	R	R 1
954-04-13	aaronha01	1976-10-03	aaroh101	2021-01-22				
## 3	Atlanta	Tommie	Aaron	Tommie Lee	190	75	R	R 1
962-04-10	aaronto01	1971-09-26	aarot101	1984-08-16				
## 4	<NA>	Don	Aase	Donald William	190	75	R	R 1
977-07-26	aasedo01	1990-10-03	aased001	<NA>				
## 5	<NA>	Andy	Abad	Fausto Andres	184	73	L	L 2
001-09-10	abadan01	2006-04-13	abada001	<NA>				
## 6	<NA>	Fernando	Abad	Fernando Antonio	235	74	L	L 2
010-07-28	abadfe01	2021-10-01	abadf001	<NA>				
## 7	Pemberton	John	Abadie	John W.	192	72	R	R 1
875-04-26	abadijo01	1875-06-10	abadj101	1905-05-17				
## 8	Fort Lauderdale	Ed	Abbatichio	Edward James	170	71	R	R 1
897-09-04	abbated01	1910-09-15	abbae101	1957-01-06				
## 9	Colchester	Bert	Abbey	Bert Wood	175	71	R	R 1
892-06-14	abbeybe01	1896-09-23	abbeb101	1962-06-11				
## 10	San Francisco	Charlie	Abbey	Charles S.	169	68	L	L 1
893-08-16	abbeych01	1897-08-19	abbec101	1926-04-27				
##	birthDate							
## 1	1981-12-27							
## 2	1934-02-05							
## 3	1939-08-05							
## 4	1954-09-08							
## 5	1972-08-25							
## 6	1985-12-17							
## 7	1850-11-04							

```
## 8 1877-04-15
## 9 1869-11-11
## 10 1866-10-14
```

The primary key in the People table is the playerID column, as it uniquely identifies each row.

b. Explain why the pair of columns { nameFirst, nameLast } aren't a key for the People table. Give an example of specific entries in the table that support your explanation.

The pair of columns { nameFirst, nameLast } are not a primary key for the People table because multiple people can have the same first and last name.

```
# find all people with the same first and last name
People %>% group_by(nameFirst, nameLast) %>% summarize(n = n()) %>% filter(n > 1)
```

```
## `summarise()` has grouped output by 'nameFirst'. You can override using the `.groups`
argument.
```

```
## # A tibble: 568 × 3
## # Groups:   nameFirst [166]
##   nameFirst nameLast     n
##   <chr>      <chr>   <int>
## 1 Abraham   Nunez         2
## 2 Adam      Eaton         2
## 3 Adam      Peterson      2
## 4 Al        Martin         2
## 5 Al        Shaw           2
## 6 Al        Smith          3
## 7 Al        Wright         2
## 8 Alberto   Castillo       2
## 9 Alex      Gonzalez       2
## 10 Alex     Sanchez        2
## # i 558 more rows
```

c. Is the column you identified in part (a) a primary key in the Batting table? Explain why or why not.

```
head(Batting, 10)
```

##	playerID	yearID	stint	teamID	lgID	G	AB	R	H	X2B	X3B	HR	RBI	SB	CS	BB	SO	IBB	H
BP	SH	SF	GIDP																
## 1	aardsda01	2004	1	SFN	NL	11	0	0	0	0	0	0	0	0	0	0	0	0	0
0 0 0	0																		
## 2	aardsda01	2006	1	CHN	NL	45	2	0	0	0	0	0	0	0	0	0	0	0	0
0 1 0	0																		
## 3	aardsda01	2007	1	CHA	AL	25	0	0	0	0	0	0	0	0	0	0	0	0	0
0 0 0	0																		
## 4	aardsda01	2008	1	BOS	AL	47	1	0	0	0	0	0	0	0	0	0	0	1	0
0 0 0	0																		
## 5	aardsda01	2009	1	SEA	AL	73	0	0	0	0	0	0	0	0	0	0	0	0	0
0 0 0	0																		
## 6	aardsda01	2010	1	SEA	AL	53	0	0	0	0	0	0	0	0	0	0	0	0	0
0 0 0	0																		
## 7	aardsda01	2012	1	NYA	AL	1	0	0	0	0	0	0	0	0	0	0	0	0	0
0 0 0	0																		
## 8	aardsda01	2013	1	NYN	NL	43	0	0	0	0	0	0	0	0	0	0	0	0	0
0 0 0	0																		
## 9	aardsda01	2015	1	ATL	NL	33	1	0	0	0	0	0	0	0	0	0	0	1	0
0 0 0	0																		
## 10	aaronha01	1954	1	ML1	NL	122	468	58	131	27	6	13	69	2	2	28	39	NA	
3 6 4	13																		

No, the playerID column is not a primary key in the Batting table because players have multiple rows in the Batting table, with each row representing a different season or year of the player’s career.

d. Is the column you identified in part (a) a foreign key in the Batting table? Explain why or why not.

Yes, the playerID column is a foreign key in the Batting table because it is a primary key in the People table and it is used to link the Batting table to the People table.

Question 4 (6 points)

Consider the following table of bird sightings; more information about this data is available at <https://github.com/rfordatascience/tidytuesday/blob/master/data/2023/2023-01-10/readme.md> (https://github.com/rfordatascience/tidytuesday/blob/master/data/2023/2023-01-10/readme.md).

```
## Rows: 100000 Columns: 22
## — Column specification —————
## Delimiter: ","
## chr (8): loc_id, subnational1_code, entry_technique, sub_id, obs_id, PROJ_PERIOD_ID,
species_code, Data_Entry_Method
## dbl (14): latitude, longitude, Month, Day, Year, how_many, valid, reviewed, day1_am,
day1_pm, day2_am, day2_pm, effort_hrs_atleast, snow_dep_a...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

a. Explain in your own words what the `distinct()` function does, and use it to explain how you know the `bird_observations` table doesn't have any duplicate rows.

The `distinct()` function is used to return a tibble with duplicate rows removed. It is used to find the unique rows in a tibble.

```
original_row_count <- nrow(bird_observations)
distinct_row_count <- nrow(distinct(bird_observations))
original_row_count
```

```
## [1] 100000
```

```
distinct_row_count
```

```
## [1] 100000
```

```
original_row_count == distinct_row_count
```

```
## [1] TRUE
```

Since the original row count is equal to the distinct row count, the `bird_observations` table doesn't have any duplicate rows.

b. Because the `bird_observations` table doesn't have any duplicate rows, it must have at least one key. What do you think is the best set of columns to choose to serve as a primary key for this table? Explain how you know it is a valid key. Hint: Think about what the observations are; you should not have more than 5 columns in your key. There is more than one correct answer.

These columns together can form a composite key because: loc_id ensures the uniqueness of the location. sub_id and obs_id ensure the uniqueness of the checklist and observation within that location. species_code specifies the species being observed. proj_period_id provides the temporal context of the observation.

We can check if the primary key is unique by using the nrow() and n_distinct() functions.

```
# Create a primary key
bird_observations <- bird_observations %>%
  mutate(primary_key = paste(loc_id, sub_id, obs_id, species_code, PROJ_PERIOD_ID, sep =
    "_"))

# Check if the primary key is unique
is_unique <- nrow(bird_observations) == n_distinct(bird_observations$primary_key)
is_unique
```

```
## [1] TRUE
```

Question 5 (6 points)

a. **Explain why the diamonds data set doesn't meet the three assumptions we discussed in class on 9-30; be specific about which assumption(s) it violates.**

The assumptions are:

1. Data is tidy
2. Order of rows does not matter
3. There are no identical rows

Lets examine the diamonds data set:

```
head(diamonds, 10)
```

```
## # A tibble: 10 × 11
##   carat cut      color clarity depth table price      x      y      z primary_key
##   <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>    <int>
## 1  0.23 Ideal      E      SI2     61.5    55   326  3.95  3.98  2.43         1
## 2  0.21 Premium    E      SI1     59.8    61   326  3.89  3.84  2.31         2
## 3  0.23 Good       E      VS1     56.9    65   327  4.05  4.07  2.31         3
## 4  0.29 Premium    I      VS2     62.4    58   334  4.2   4.23  2.63         4
## 5  0.31 Good       J      SI2     63.3    58   335  4.34  4.35  2.75         5
## 6  0.24 Very Good  J      VVS2    62.8    57   336  3.94  3.96  2.48         6
## 7  0.24 Very Good  I      VVS1    62.3    57   336  3.95  3.98  2.47         7
## 8  0.26 Very Good  H      SI1     61.9    55   337  4.07  4.11  2.53         8
## 9  0.22 Fair       E      VS2     65.1    61   337  3.87  3.78  2.49         9
## 10 0.23 Very Good  H      VS1     59.4    61   338  4     4.05  2.39        10
```

```
# Check if the order of rows matters
order_column_exists <- any(names(diamonds) %in% "order")
order_column_exists
```

```
## [1] FALSE
```

```
# Check for duplicate rows
duplicate_rows_exist <- nrow(diamonds) != nrow(distinct(diamonds))
duplicate_rows_exist
```

```
## [1] FALSE
```

```
# Identify duplicate rows
duplicate_rows <- diamonds %>%
  group_by(across(everything())) %>%
  filter(n() > 1) %>%
  ungroup()
```

```
# Display the duplicate rows
duplicate_rows
```

```
## # A tibble: 0 × 11
## # i 11 variables: carat <dbl>, cut <ord>, color <ord>, clarity <ord>, depth <dbl>, ta
ble <dbl>, price <int>, x <dbl>, y <dbl>, z <dbl>,
## #   primary_key <int>
```

The diamonds data set violates the third assumption because there are duplicate rows.

b. Suppose you have good documentation for the diamonds data set and know that: every row corresponds to a different diamond; there were no mistakes in collecting/entering this data; and the order of the diamonds doesn't matter. Modify the diamonds data set in an appropriate way to make sure it satisfies the three assumptions.

I will add a primary key to the diamonds data set.

```
diamonds <- diamonds %>%
  mutate(primary_key = row_number())
```


Question 6 (21 points)

Consider the following tibbles:

- a. Suppose the order of the cities in CityInfo list matters. Modify your table in an appropriate way, and explain why this is a good thing to do.

I will add a primary key to the CityInfo table. This is a good thing to do because it ensures that the order of the cities in the CityInfo table is preserved.

```
CityInfo <- CityInfo %>%
  mutate(primary_key = row_number())

head(CityInfo, 10)
```

```
## # A tibble: 8 × 4
##   City          Country Population primary_key
##   <chr>         <chr>      <dbl>      <int>
## 1 Boston        USA          650706          1
## 2 San Jose      Costa Rica    339581          2
## 3 Toronto       Canada     2930000          3
## 4 Rio de Janeiro Brazil     6211000          4
## 5 Cartago       Costa Rica    160457          5
## 6 Vancouver     Canada     675218          6
## 7 Buenos Aires  Argentina   3121000          7
## 8 Los Angeles   USA        3822000          8
```

- b. Join these tibbles according to Country using an inner join. Which city/cities appear in two different rows, which city/cities appear only in one row, and which city/cities don't appear in this tibble? Explain why this is.

```
inner_join(CityInfo, Regions, by = "Country")
```

```
## Warning in inner_join(CityInfo, Regions, by = "Country"): Detected an unexpected many
-to-many relationship between `x` and `y`.
## i Row 2 of `x` matches multiple rows in `y`.
## i Row 2 of `y` matches multiple rows in `x`.
## i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to
silence this warning.
```

```
## # A tibble: 9 × 5
##   City          Country      Population primary_key Region
##   <chr>         <chr>         <dbl>         <int> <chr>
## 1 Boston      USA             650706         1 North America
## 2 San Jose    Costa Rica      339581         2 Central America
## 3 San Jose    Costa Rica      339581         2 North America
## 4 Toronto     Canada          2930000        3 North America
## 5 Rio de Janeiro Brazil          6211000        4 South America
## 6 Cartago     Costa Rica      160457         5 Central America
## 7 Cartago     Costa Rica      160457         5 North America
## 8 Vancouver   Canada          675218         6 North America
## 9 Los Angeles USA             3822000        8 North America
```

Costa Rica appears in two different rows because it is a country that is both in Central America and North America. The cities that appear in two different rows are San Jose and Cartago. The cities that appear only in one row are Boston, Toronto, Rio de Janeiro, Vancouver, and Los Angeles. The cities that don't appear in this tibble are Panama City and Santiago. This is because Panama and Chile are not listed in the Regions tibble.

c. Joining these tables with a `left_join` rather than an `inner_join` results in a tibble with one more row than in part (b). Which additional row is present here and why?

```
left_join(CityInfo, Regions, by = "Country")
```

```
## Warning in left_join(CityInfo, Regions, by = "Country"): Detected an unexpected many-
to-many relationship between `x` and `y`.
## i Row 2 of `x` matches multiple rows in `y`.
## i Row 2 of `y` matches multiple rows in `x`.
## i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to
silence this warning.
```

```
## # A tibble: 10 × 5
##   City          Country      Population primary_key Region
##   <chr>         <chr>         <dbl>         <int> <chr>
## 1 Boston      USA             650706         1 North America
## 2 San Jose    Costa Rica      339581         2 Central America
## 3 San Jose    Costa Rica      339581         2 North America
## 4 Toronto     Canada          2930000        3 North America
## 5 Rio de Janeiro Brazil          6211000        4 South America
## 6 Cartago     Costa Rica      160457         5 Central America
## 7 Cartago     Costa Rica      160457         5 North America
## 8 Vancouver   Canada          675218         6 North America
## 9 Buenos Aires Argentina       3121000        7 <NA>
## 10 Los Angeles USA             3822000        8 North America
```

The additional row is Argentina. This is because Argentina is a country that is only in the CityInfo table.

d. **Joining these tables with a `right_join` rather than an `inner_join` results in a tibble with two more rows than in part (b). Which additional rows are present here and why?**

```
right_join(CityInfo, Regions, by = "Country")
```

```
## Warning in right_join(CityInfo, Regions, by = "Country"): Detected an unexpected many-
to-many relationship between `x` and `y`.
## i Row 2 of `x` matches multiple rows in `y`.
## i Row 2 of `y` matches multiple rows in `x`.
## i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to
silence this warning.
```

```
## # A tibble: 11 × 5
##   City          Country      Population primary_key Region
##   <chr>         <chr>          <dbl>      <int> <chr>
## 1 Boston      USA             650706         1 North America
## 2 San Jose    Costa Rica      339581         2 Central America
## 3 San Jose    Costa Rica      339581         2 North America
## 4 Toronto     Canada          2930000        3 North America
## 5 Rio de Janeiro Brazil          6211000        4 South America
## 6 Cartago     Costa Rica      160457         5 Central America
## 7 Cartago     Costa Rica      160457         5 North America
## 8 Vancouver   Canada          675218         6 North America
## 9 Los Angeles USA             3822000        8 North America
## 10 <NA>        Panama          NA          NA Central America
## 11 <NA>        Chile           NA          NA South America
```

The additional rows are Panama and Chile. This is because Panama and Chile are countries that are only in the Regions tibble.

e. **Joining these tables with a `full_join` rather than an `inner_join` results in a tibble with three more rows than in part (b). Which additional rows are present here and why?**

```
full_join(CityInfo, Regions, by = "Country")
```

```
## Warning in full_join(CityInfo, Regions, by = "Country"): Detected an unexpected many-
to-many relationship between `x` and `y`.
## i Row 2 of `x` matches multiple rows in `y`.
## i Row 2 of `y` matches multiple rows in `x`.
## i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to
silence this warning.
```

```
## # A tibble: 12 × 5
##   City          Country      Population primary_key Region
##   <chr>         <chr>          <dbl>         <int> <chr>
## 1 Boston        USA             650706          1 North America
## 2 San Jose      Costa Rica      339581          2 Central America
## 3 San Jose      Costa Rica      339581          2 North America
## 4 Toronto       Canada          2930000         3 North America
## 5 Rio de Janeiro Brazil          6211000         4 South America
## 6 Cartago       Costa Rica      160457          5 Central America
## 7 Cartago       Costa Rica      160457          5 North America
## 8 Vancouver     Canada          675218          6 North America
## 9 Buenos Aires Argentina       3121000         7 <NA>
## 10 Los Angeles  USA             3822000         8 North America
## 11 <NA>         Panama          NA             NA Central America
## 12 <NA>         Chile           NA             NA South America
```

The additional rows are Panama, Chile, and Argentina. This is because Panama and Chile are countries that are only in the Regions tibble, and Argentina is a country that is only in the CityInfo table.

f. Join these tibbles according to Country using a `semi_join()`. Which row(s) and column(s) appear in the resulting table? Explain why this is.

```
semi_join(CityInfo, Regions, by = "Country")
```

```
## # A tibble: 7 × 5
##   City          Country      Population primary_key
##   <chr>         <chr>          <dbl>         <int>
## 1 Boston        USA             650706          1
## 2 San Jose      Costa Rica      339581          2
## 3 Toronto       Canada          2930000         3
## 4 Rio de Janeiro Brazil          6211000         4
## 5 Cartago       Costa Rica      160457          5
## 6 Vancouver     Canada          675218          6
## 7 Los Angeles  USA             3822000         8
```

The rows that appear in the resulting table are those from CityInfo where the Country also appears in the Regions table. The columns in the resulting table are the same as those in the CityInfo table. This is because `semi_join()` filters CityInfo to only include rows with a Country that is present in Regions.

g. Join these tibbles according to Country using an `anti_join()`. Which row(s) and column(s) appear in the resulting table? Explain why this is.

```
anti_join(CityInfo, Regions, by = "Country")
```

```
## # A tibble: 1 × 4
##   City      Country Population primary_key
##   <chr>    <chr>      <dbl>      <int>
## 1 Buenos Aires Argentina    3121000         7
```

The rows that appear in the resulting table are those from CityInfo where the Country does not appear in the Regions table. The columns in the resulting table are the same as those in the CityInfo table. This is because `anti_join()` filters CityInfo to exclude rows with a Country that is present in Regions.

Question 7 (6 points)

Consider the following two tibbles.

a. Join these tibbles by the species column using a `full_join`. Explain why doing this join is probably a bad idea.

```
full_join(October_Pets, Pet_Average_Weights, by = "species")
```

```
## Warning in full_join(October_Pets, Pet_Average_Weights, by = "species"): Detected an
## unexpected many-to-many relationship between `x` and `y`.
## i Row 1 of `x` matches multiple rows in `y`.
## i Row 3 of `y` matches multiple rows in `x`.
## i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to
## silence this warning.
```

```
## # A tibble: 20 × 6
##   name      species age_months arrival_day sex      avg_weight_lbs
##   <chr>    <chr>      <dbl>      <dbl> <chr>      <dbl>
## 1 Sparky   Dog          31          3 Female      45
## 2 Sparky   Dog          31          3 Male       50
## 3 Fido     Dog          29         11 Female      45
## 4 Fido     Dog          29         11 Male       50
## 5 Fluffy   Cat          78          4 Female      9.4
## 6 Fluffy   Cat          78          4 Male     10.1
## 7 Lassie   Dog          98         28 Female      45
## 8 Lassie   Dog          98         28 Male       50
## 9 Patches  Cat         115         14 Female      9.4
## 10 Patches  Cat         115         14 Male     10.1
## 11 Spot     Dog           7         12 Female      45
## 12 Spot     Dog           7         12 Male       50
## 13 Socks    Cat           4         17 Female      9.4
## 14 Socks    Cat           4         17 Male     10.1
## 15 Buddy    Dog          15         15 Female      45
## 16 Buddy    Dog          15         15 Male       50
## 17 Lizzie   Lizard        2           1 Female      0.4
## 18 Lizzie   Lizard        2           1 Male       0.3
## 19 Tweety   Bird           6           2 Female      0.8
## 20 Tweety   Bird           6           2 Male       0.9
```

This join is probably a bad idea because it is a many-to-many relationship. This is because there are multiple species in the October_Pets table that have the same species in the Pet_Average_Weights table.

b. Explain why you have the number of rows that you do in your join in the previous part.

The number of rows in the join is 20 because each row in the October_Pets table is matched with each corresponding row in the Pet_Average_Weights table based on the species column. This results in a many-to-many relationship due to the gender differences, leading to multiple rows for each species and gender combination.

Question 8 (9 points)

a. (6 points) Add to the flights data set the altitude of the origin airports and the altitude of the destination airports. That is, each row should now have 2 more additional columns, which you should name origin_alt and dest_alt. Move your columns for origin, destination, and their altitudes to the front of your data set, with the remaining columns displayed after them.

```

flights <- flights %>%
  left_join(select(airports, faa, alt), by = c("origin" = "faa")) %>%
  left_join(select(airports, faa, alt), by = c("dest" = "faa"), suffix = c("_origin", "_dest"))

flights <- flights %>%
  select(origin, dest, alt_origin, alt_dest, everything())

head(flights, 10)

```

```

## # A tibble: 10 × 29
##   origin dest alt_origin alt_dest year month day dep_time sched_dep_time dep_del
ay arr_time sched_arr_time arr_delay carrier flight tailnum
##   <chr> <chr> <dbl> <dbl> <int> <int> <int> <int> <int> <db
l> <int> <int> <int> <dbl> <chr> <int> <chr>
## 1 EWR IAH 18 97 2013 1 1 517 515
2 830 819 11 UA 1545 N14228
## 2 LGA IAH 22 97 2013 1 1 533 529
4 850 830 20 UA 1714 N24211
## 3 JFK MIA 13 8 2013 1 1 542 540
2 923 850 33 AA 1141 N619AA
## 4 JFK BQN 13 NA 2013 1 1 544 545
-1 1004 1022 -18 B6 725 N804JB
## 5 LGA ATL 22 1026 2013 1 1 554 600
-6 812 837 -25 DL 461 N668DN
## 6 EWR ORD 18 668 2013 1 1 554 558
-4 740 728 12 UA 1696 N39463
## 7 EWR FLL 18 9 2013 1 1 555 600
-5 913 854 19 B6 507 N516JB
## 8 LGA IAD 22 313 2013 1 1 557 600
-3 709 723 -14 EV 5708 N829AS
## 9 JFK MCO 13 96 2013 1 1 557 600
-3 838 846 -8 B6 79 N593JB
## 10 LGA ORD 22 668 2013 1 1 558 600
-2 753 745 8 AA 301 N3ALAA
## # i 13 more variables: air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time
_hour <dtm>, alt_origin_origin <dbl>, alt_dest_dest <dbl>,
## # alt_origin_origin_origin <dbl>, alt_dest_dest_dest <dbl>, alt_origin_origin_origi
n_origin <dbl>, alt_dest_dest_dest_dest <dbl>,
## # alt_origin_origin_origin_origin_origin <dbl>, alt_dest_dest_dest_dest_dest <dbl>

```

b. (3 points) **The following command attaches plane information to the flights tibble, for all flights where the tail number appears in the planes tibble. There's over 284,000 such flights:**

```
inner_join(flights, planes, by = "tailnum")
```

```
## # A tibble: 284,170 × 37
##   origin dest alt_origin alt_dest year.x month   day dep_time sched_dep_time dep_de
lay arr_time sched_arr_time arr_delay carrier flight tailnum
##   <chr>  <chr>      <dbl>    <dbl> <int> <int> <int>    <int>      <int>    <d
bl>    <int>      <int>      <dbl> <chr>    <int> <chr>
## 1 EWR    IAH        18      97   2013    1    1      517        515
2      830      819      11 UA      1545 N14228
## 2 LGA    IAH        22      97   2013    1    1      533        529
4      850      830      20 UA      1714 N24211
## 3 JFK    MIA        13       8   2013    1    1      542        540
2      923      850      33 AA      1141 N619AA
## 4 JFK    BQN        13      NA   2013    1    1      544        545
-1     1004     1022     -18 B6      725 N804JB
## 5 LGA    ATL        22     1026  2013    1    1      554        600
-6      812      837     -25 DL      461 N668DN
## 6 EWR    ORD        18     668   2013    1    1      554        558
-4      740      728     12 UA      1696 N39463
## 7 EWR    FLL        18       9   2013    1    1      555        600
-5      913      854     19 B6      507 N516JB
## 8 LGA    IAD        22     313   2013    1    1      557        600
-3      709      723     -14 EV      5708 N829AS
## 9 JFK    MCO        13      96   2013    1    1      557        600
-3      838      846     -8 B6      79 N593JB
## 10 JFK   PBI        13      19   2013    1    1      558        600
-2      849      851     -2 B6      49 N793JB
## # i 284,160 more rows
## # i 21 more variables: air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time
_hour <dtm>, alt_origin_origin <dbl>, alt_dest_dest <dbl>,
## #   alt_origin_origin_origin <dbl>, alt_dest_dest_dest <dbl>, alt_origin_origin_origi
n_origin <dbl>, alt_dest_dest_dest_dest <dbl>,
## #   alt_origin_origin_origin_origin_origin <dbl>, alt_dest_dest_dest_dest_dest <dbl>,
year.y <int>, type <chr>, manufacturer <chr>, model <chr>,
## #   engines <int>, seats <int>, speed <int>, engine <chr>
```

When we remove the “by” argument, we get a tibble with fewer than 5000 rows. Explain what’s happening here, and why these particular rows have been included in this tibble.

```
inner_join(flights, planes)
```

```
## Joining with `by = join_by(year, tailnum)`
```



```
## # A tibble: 4,630 × 36
##   origin dest alt_origin alt_dest year month day dep_time sched_dep_time dep_delay
##   <chr> <chr> <dbl> <dbl> <int> <int> <int> <int> <int> <int>
##   1 EWR FLL 18 9 2013 1 18 1846 1810
##   2 JFK BOS 13 19 2013 10 1 647 655
##   3 EWR LAX 18 126 2013 10 1 652 652
##   4 JFK MSP 13 841 2013 10 1 755 800
##   5 JFK HOU 13 46 2013 10 1 813 820
##   6 JFK SYR 13 421 2013 10 1 925 930
##   7 JFK IAD 13 313 2013 10 1 1113 1120
##   8 JFK ROC 13 559 2013 10 1 1426 1429
##   9 LGA CLT 22 748 2013 10 1 1446 1450
##  10 JFK MSY 13 4 2013 10 1 1454 1455
## # i 4,620 more rows
## # i 20 more variables: air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time
## # alt_origin_origin_origin <dbl>, alt_dest_dest_dest <dbl>, alt_origin_origin_origin
## # alt_origin_origin_origin_origin_origin <dbl>, alt_dest_dest_dest_dest_dest <dbl>,
## # type <chr>, manufacturer <chr>, model <chr>, engines <int>,
## # seats <int>, speed <int>, engine <chr>
```

When we remove the “by” argument, we get a tibble with fewer than 5000 rows because the join is performed using all columns that have the same names in both tibbles. The rows that are included in this tibble are the ones where all corresponding columns in both tibbles match.

Question 9 (18 points)

This question considers the following three data sets, from a sentiment analysis for African languages. More information about this data set can be found at <https://github.com/rfordatascience/tidytuesday/blob/master/data/2023/2023-02-28/readme.md> (<https://github.com/rfordatascience/tidytuesday/blob/master/data/2023/2023-02-28/readme.md>).

```
afrisenti <- read_csv("https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2023/2023-02-28/afrisenti.csv")
```

```
## Rows: 111720 Columns: 4
```

```
## — Column specification —————
```

```
## Delimiter: ","
```

```
## chr (4): language_iso_code, tweet, label, intended_use
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
languages <- read_csv("https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2023/2023-02-28/languages.csv")
```

```
## Rows: 14 Columns: 2
```

```
## — Column specification —————
```

```
## Delimiter: ","
```

```
## chr (2): language_iso_code, language
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
language_countries <- read_csv("https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2023/2023-02-28/language_countries.csv")
```

```
## Rows: 23 Columns: 2
```

```
## — Column specification —————
```

```
## Delimiter: ","
```

```
## chr (2): language_iso_code, country
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

a. Are there any language iso codes that appear in the afrisenti table but not in the languages table, or vice versa? Explain how you know.

```
missing_in_languages <- afrisenti %>%  
  anti_join(languages, by = "language_iso_code")  
missing_in_languages
```

```
## # A tibble: 0 × 4
## # i 4 variables: language_iso_code <chr>, tweet <chr>, label <chr>, intended_use <chr>
>
```

```
missing_in_afrisenti <- languages %>%
  anti_join(afrisenti, by = "language_iso_code")
missing_in_afrisenti
```

```
## # A tibble: 0 × 2
## # i 2 variables: language_iso_code <chr>, language <chr>
```

No, there are no language iso codes that appear in the afrisenti table but not in the languages table, or vice versa because the `anti_join()` function returns an empty tibble.

b. Explain why a `left_join`, `right_join`, `inner_join`, and `full_join` of the `afrisenti` and `languages` data tables will all produce the same result.

A `left_join`, `right_join`, `inner_join`, and `full_join` of the `afrisenti` and `languages` tables will all produce the same result because the language iso codes match in both tables.

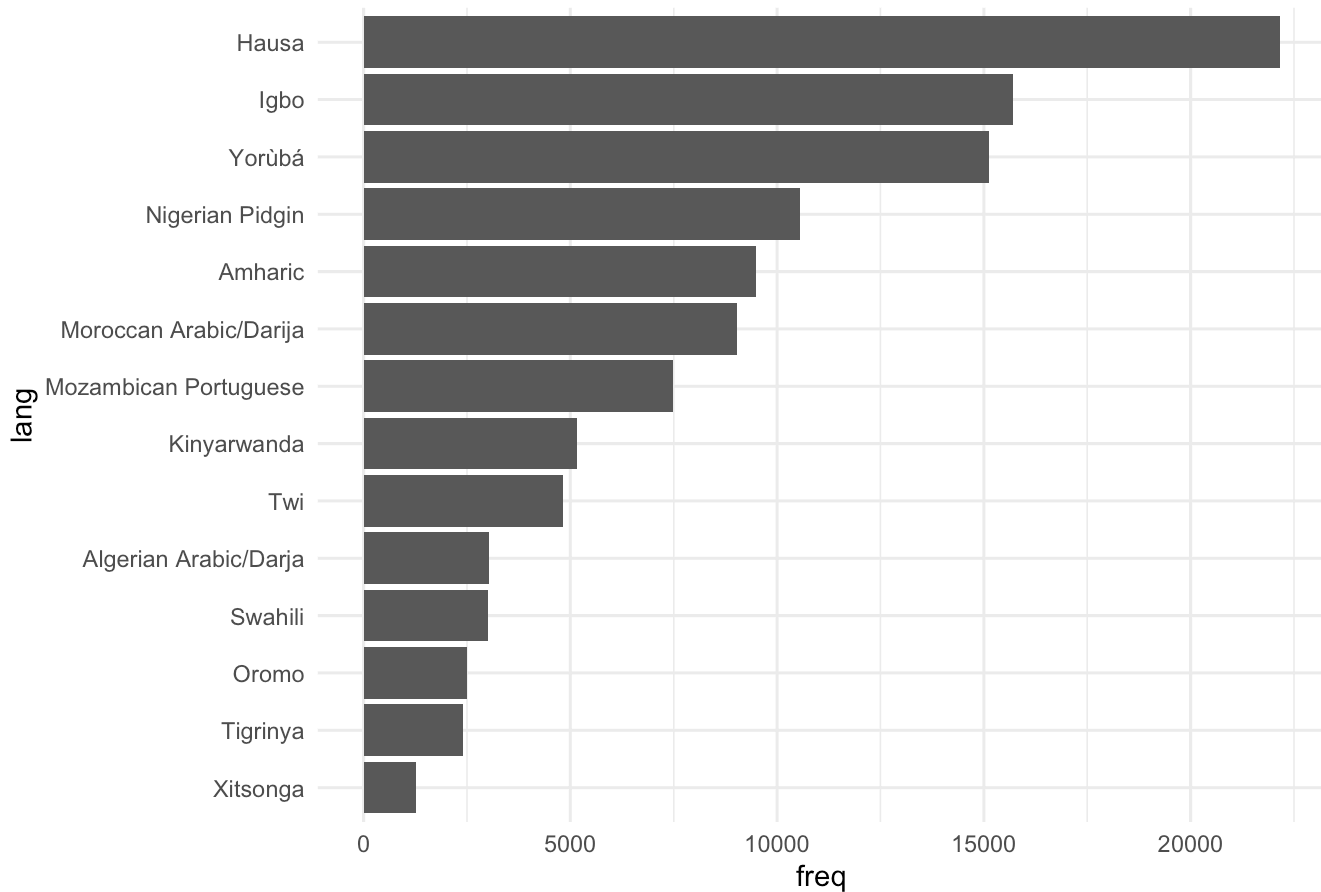
c. Make a barchart showing how frequently each language appears in the `afrisenti` table. Your plot should use the full name of all the languages, not the iso abbreviations for the languages.

```
afrisenti_full <- afrisenti %>%
  left_join(languages, by = "language_iso_code")

language_counts <- afrisenti_full %>%
  count(language)

ggplot(language_counts, aes(x = reorder(language, n), y = n)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  labs(
    title = "Frequency of Each Language in Afrisenti Table",
    x = "lang",
    y = "freq"
  ) +
  theme_minimal()
```

Frequency of Each Language in Afrisenti Table



d. Join the `afrisenti` and `language_countries` data sets using any join type you'd like. Explain why, although the `afrisenti` table has 111,720 rows, the new joined table has 186,941 rows. *Note: Some systems may have trouble knitting the 'tweet' column due to the special characters present, so if this applies to you, feel free to add `%>% select(-tweet)` to your answer.*

```
inner_join(afrisenti %>% select(-tweet), language_countries)
```

```
## Joining with `by = join_by(language_iso_code)`
```

```
## Warning in inner_join(afrisenti %>% select(-tweet), language_countries): Detected an
unexpected many-to-many relationship between `x` and `y`.
## i Row 21542 of `x` matches multiple rows in `y`.
## i Row 1 of `y` matches multiple rows in `x`.
## i If a many-to-many relationship is expected, set `relationship = "many-to-many"` to
silence this warning.
```

```
## # A tibble: 186,941 × 4
##   language_iso_code label      intended_use country
##   <chr>              <chr>      <chr>      <chr>
## 1 amh                negative dev      Ethiopia
## 2 amh                negative dev      Ethiopia
## 3 amh                negative dev      Ethiopia
## 4 amh                negative dev      Ethiopia
## 5 amh                negative dev      Ethiopia
## 6 amh                negative dev      Ethiopia
## 7 amh                negative dev      Ethiopia
## 8 amh                negative dev      Ethiopia
## 9 amh                negative dev      Ethiopia
## 10 amh               negative dev      Ethiopia
## # i 186,931 more rows
```

The increase in the number of rows from 111,720 to 186,941 is due to the many-to-many relationship between the `afrisenti` and `language_countries` datasets. Each language in `afrisenti` can be associated with multiple countries in `language_countries`, leading to multiple rows in the joined dataset for each original row in `afrisenti`.

e. Make a table consisting only of the 8 languages appearing most frequently in the `afrisenti` table. Your table should only have 8 rows, one for each of these languages.

```
afrisenti_full <- afrisenti %>%
  left_join(languages, by = "language_iso_code")

language_counts <- afrisenti_full %>%
  count(language)

language_counts %>%
  arrange(desc(n)) %>%
  head(8)
```

```
## # A tibble: 8 × 2
##   language      n
##   <chr>      <int>
## 1 Hausa      22152
## 2 Igbo       15715
## 3 Yorùbá     15127
## 4 Nigerian Pidgin 10556
## 5 Amharic      9480
## 6 Moroccan Arabic/Darija 9038
## 7 Mozambican Portuguese 7492
## 8 Kinyarwanda  5155
```

f. Filter the `afrisenti` table, using a join we learned this week, to only keep rows corresponding to one of the 8 languages that appears most frequently in the table. Hint: Use your table from the previous part. *Note: Some systems may have trouble knitting the 'tweet' column due to the special characters present, so feel free to add `%>% select(-tweet)` to your answer.*

```
language_counts <- afrisenti %>%
  count(language_iso_code, sort = TRUE)

top_languages <- language_counts %>%
  top_n(8, n) %>%
  select(language_iso_code)

filtered_afrisenti <- afrisenti %>%
  semi_join(top_languages, by = "language_iso_code") %>%
  select(-tweet)

filtered_afrisenti
```

```
## # A tibble: 94,715 × 3
##   language_iso_code label      intended_use
##   <chr>             <chr>      <chr>
## 1 amh               negative dev
## 2 amh               negative dev
## 3 amh               negative dev
## 4 amh               negative dev
## 5 amh               negative dev
## 6 amh               negative dev
## 7 amh               negative dev
## 8 amh               negative dev
## 9 amh               negative dev
## 10 amh              negative dev
## # i 94,705 more rows
```

Question 10 (6 points)

Filter the flights data set to only contain flights along the 20 routes with the largest average arrival delays (of the flights that took off), where a route consists of both the origin airport and the destination airport. Hint: You may want to make an intermediate table to help you.

```
flights_with_avg_delay <- flights %>%
  group_by(origin, dest) %>%
  summarise(avg_arr_delay = mean(arr_delay, na.rm = TRUE)) %>%
  ungroup()
```

`summarise()` has grouped output by 'origin'. You can override using the `.groups` argument.

```
top_routes <- flights_with_avg_delay %>%
  arrange(desc(avg_arr_delay)) %>%
  slice_max(order_by = avg_arr_delay, n = 20)

filtered_flights <- flights %>%
  semi_join(top_routes, by = c("origin", "dest"))

head(filtered_flights)
```

```
## # A tibble: 6 × 29
##   origin dest alt_origin alt_dest year month day dep_time sched_dep_time dep_delay
##   <chr> <chr> <dbl> <dbl> <int> <int> <int> <int> <int> <dbl>
##   <int> <int> <int> <dbl> <chr> <int> <chr>
## 1 EWR MEM 18 341 2013 1 1 812 814 -
## 2 1040 1017 23 EV 4537 N17108
## 2 EWR JAC 18 6451 2013 1 1 848 851 -
## 3 1155 1136 19 UA 1741 N27724
## 3 EWR DCA 18 15 2013 1 1 929 929
## 4 1028 1042 -14 EV 4636 N11551
## 4 EWR MKE 18 723 2013 1 1 1044 1045 -
## 5 1231 1212 19 EV 4322 N15555
## 5 EWR PWM 18 77 2013 1 1 1056 1059 -
## 6 1203 1209 -6 EV 4479 N11544
## 6 LGA CAK 22 1228 2013 1 1 1147 1155 -
## 8 1335 1327 8 FL 353 N932AT
## # i 13 more variables: air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time
## # alt_origin_origin <dbl>, alt_dest_dest <dbl>,
## # alt_origin_origin_origin <dbl>, alt_dest_dest_dest <dbl>, alt_origin_origin_ori
## # alt_dest_dest_dest_dest <dbl>,
## # alt_origin_origin_origin_origin <dbl>, alt_dest_dest_dest_dest_dest <dbl>
```

Question 11 (6 points)

- Explain what R's `intersect()` function is, explain how it is different from an `inner_join`, and give a real-world example of when you might want to use it.

The `intersect()` function returns common rows between two tables, unlike `inner_join` which merges columns. An example of when you might want to use it is when you have two tables of voters and you want to find the voters that are present in both tables.

b. Explain what R's `setdiff()` function is, explain how it is different from an `anti_join`, and give a real-world example of when you might want to use it.

The `setdiff()` function returns rows in one table that are not present in another table, unlike `anti_join` which filters rows. An example of when it would be useful is when you have two tables of trinkets and you want to find the trinkets that are present in one table but not in the other.

Question 12 (6 points)

Consider the following three data sets with more detailed information about three of the African languages considered above. More information about this data can be found at https://github.com/afrisenti-semeval/afrisenti-semeval-2023/tree/main/data_with_annotators_labels#readme (https://github.com/afrisenti-semeval/afrisenti-semeval-2023/tree/main/data_with_annotators_labels#readme).

```
morrocan_arabic <- read_csv("https://raw.githubusercontent.com/afrisenti-semeval/afrisenti-semeval-2023/main/data_with_annotators_labels/morrocan_arabic_individual_labels.csv")
```

```
## Rows: 6999 Columns: 8
```

```
## — Column specification
```

```
## Delimiter: ","
```

```
## chr (4): text, label_1, label_2, label_3
```

```
## dbl (4): text_id, annotator_1, annotator_2, annotator_3
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
algerian_arabic <- read_csv("https://raw.githubusercontent.com/afrisenti-semeval/afrisenti-semeval-2023/main/data_with_annotators_labels/algerian_arabic_individual_labels.csv")
```



```
## Rows: 3097 Columns: 4
## — Column specification —————
## Delimiter: ","
## chr (3): annotator1, annotator2, annotator3
## dbl (1): tweet_id
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
hausan <- read_csv("https://raw.githubusercontent.com/afrisenti-semeval/afrisent-semeval-2023/main/data_with_annotators_labels/hausan_individual_labels.csv")
```

```
## Rows: 30000 Columns: 8
## — Column specification —————
## Delimiter: ","
## chr (4): text, label_1, label_2, label_3
## dbl (4): text_id, annotator_1, annotator_2, annotator_3
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

a. Combine the morrocan_arabic and hausan data sets together into a single table in an appropriate way. You should be sure your resulting table contains information about which rows are observations about Moroccan Arabic and which rows are observations about Hausa.

```
combined_data <- bind_rows(
  morrocan_arabic %>% mutate(language = "Morrocan Arabic"),
  hausan %>% mutate(language = "Hausa")
)

combined_data
```

```
## # A tibble: 36,999 × 9
##   text                                                                 text_id annotat
or_1 annotator_2 annotator_3 label_1 label_2 label_3 language
##   <chr>                                                                 <dbl>      <
dbl>      <dbl>      <dbl> <chr>   <chr>   <chr>   <chr>
## 1 "#nada0074 Jomo3a mobaraka inchallah 3liya we 3la la famille dyal... 7.18e17
79          72          73 Positi... Positi... Positi... Morroca...
## 2 "#WhatAboutArabArmy Lay lay lay lay lay lay la la la lay lay 🤔🤔 ... 1.13e18
79          72          73 Negati... Indete... Indete... Morroca...
## 3 "@nohita123 @Anyssa_Ch la daba homa li ghadi ykhtaro hna khas nsd... 5.72e17
72          79          73 Neutral Neutral Positi... Morroca...
## 4 "@sansuuna matbkhlich 3lina wakha ma3rt fin kati7o 3la had la9ata... 6.78e17
72          79          73 Neutral Neutral Positi... Morroca...
## 5 "@aminattttta o soltana fella wa3dat ibtissam boghniya"                5.99e17
73          72          79 Neutral Neutral Neutral Morroca...
## 6 "@Ihab_Amir ihaaab nta tstaleel la9ab o nta charaftii l maghreeb ... 6.83e17
72          79          73 Positi... Positi... Positi... Morroca...
## 7 "@greenadilaida @FatihiW Merehba khouya adil chi poisson au four ... 1.19e18
79          73          72 Neutral Positi... Positi... Morroca...
## 8 "Fach l prof dyal communication katsm3ek glti \"un video\" https:... 1.11e18
73          79          72 Neutral Neutral Neutral Morroca...
## 9 "@91Grosminey @IbtissamTiskat @fatiinatiskat @ali__shaddad @Fulla... 5.98e17
73          72          79 Negati... Negati... Negati... Morroca...
## 10 "@_BigBen__ @L7argouss @ucef79 kayn Chi man9diw ? Sme3t sou9 rass... 7.42e17
72          79          73 Neutral Neutral Neutral Morroca...
## # i 36,989 more rows
```

b. Explain why you can't combine the morrocan_arabic and algerian_arabic tables together in the same way you did in the previous part for morrocan_arabic and hausa.

You can't combine the morrocan_arabic and algerian_arabic tables together in the same way you did in the previous part for morrocan_arabic and hausa because the columns in the two tables are different.